



Εθνικό Μετσόβιο Πολυτεχνείο
Σχολή Ηλεκτρολόγων Μηχανικών και Μηχανικών Υπολογιστών
Τομέας Τεχνολογίας Πληροφορικής και Υπολογιστών

Control Flow Integrity

Διπλωματική Εργασία

του

Νίκου Γιανναράκη

Επιβλέπων: Κωστής Σαγώνας
Αν. Καθηγητής Ε.Μ.Π.

Εργαστήριο Τεχνολογίας Λογισμικού
Αθήνα, Σεπτέμβριος 2014



Εθνικό Μετσόβιο Πολυτεχνείο
Σχολή Ηλεκτρολόγων Μηχανικών και Μηχανικών Υπολογιστών
Τομέας Τεχνολογίας Πληροφορικής και Υπολογιστών
Εργαστήριο Τεχνολογίας Λογισμικού

Control Flow Integrity

Διπλωματική Εργασία

του

Νίκου Γιανναράκη

Επιβλέπων: Κωστής Σαγώνας
Αν. Καθηγητής Ε.Μ.Π.

Εγκρίθηκε από την τριμελή εξεταστική επιτροπή την 11^η Σεπτεμβρίου, 2014.

.....
Κωστής Σαγώνας
Αν. Καθηγητής Ε.Μ.Π.

.....
Νικόλαος Παπασπύρου
Αν. Καθηγητής Ε.Μ.Π.

.....
Ιώαννης Σμαραγδάκης
Αν. Καθηγητής Ε.Μ.Π.

Αθήνα, Σεπτέμβριος 2014

.....
Nick Giannarakis
Διπλωματούχος Ηλεκτρολόγος Μηχανικός
και Μηχανικός Υπολογιστών Ε.Μ.Π.

Copyright © – All rights reserved Nick Giannarakis, 2014.

Με επιφύλαξη παντός δικαιώματος.

Απαγορεύεται η αντιγραφή, αποθήκευση και διανομή της παρούσας εργασίας, εξ ολοκλήρου ή τμήματος αυτής, για εμπορικό σκοπό. Επιτρέπεται η ανατύπωση, αποθήκευση και διανομή για σκοπό μη κερδοσκοπικό, εκπαιδευτικής ή ερευνητικής φύσης, υπό την προϋπόθεση να αναφέρεται η πηγή προέλευσης και να διατηρείται το παρόν μήνυμα. Ερωτήματα που αφορούν τη χρήση της εργασίας για κερδοσκοπικό σκοπό πρέπει να απευθύνονται προς τον συγγραφέα.

Οι απόψεις και τα συμπεράσματα που περιέχονται σε αυτό το έγγραφο εκφράζουν τον συγγραφέα και δεν πρέπει να ερμηνευθεί ότι αντιπροσωπεύουν τις επίσημες θέσεις του Εθνικού Μετσόβιου Πολυτεχνείου.

The purpose of this diploma thesis is to present a novel, hardware-assisted, formally verified implementation of low-level security policies, such as Control-Flow Integrity and Call Stack Protection. Contrary to existing

Keywords

concolic testing, Erlang, software testing, dynamic symbolic execution, SMT solving

Ευχαριστίες

Θα ήθελα να ευχαριστήσω τον Κωστή Σαγώνα για την καθοδήγησή του, την υπομονή του και την πίστη του σε εμένα. Σε μια περίεργη περίοδο της ακαδημαϊκής μου πορείας, η εμπιστοσύνη που μου έδειξε αποτέλεσε το σημαντικότερο κίνητρό μου.

Επίσης θα ήθελα να ευχαριστήσω τον Νίκο Παπασπύρου, καταρχάς για την πολύτιμη βοήθεια του, αλλά πολύ περισσότερο για το ότι αποτελεί τον βασικό λόγο για τον οποίο ασχολήθηκα με την Πληροφορική. Η διδασκαλία του και το ήθος του ήταν και συνεχίζουν να είναι έμπνευση για εμένα.

Ευχαριστώ την οικογένειά μου που με έκανε τον άνθρωπο που είμαι σήμερα.

Τέλος, ευχαριστώ την Μαιρούλα μου για την στήριξη της όλα αυτά τα χρόνια. Σε ευχαριστώ που είσαι εσύ και που με κάνεις καλύτερο!

Άγγελος Γιάντσος

Contents

	5
Ευχαριστίες	7
Contents	10
List of Figures	11
Listings	13
1 Introduction	15
1.1 Motivation	15
1.2 Thesis Outline	16
2 Safety and Security Policies	17
2.1 A Programmable Unit for Metadata Processing	17
2.2 Micro-policies: A Framework for Verified, Hardware-Assisted Security Monitors	18
2.2.1 Correctness of micro-policies	19
2.2.2 Basic Machine	19
2.2.3 Symbolic Machine	20
2.2.4 Concrete Machine	20
2.2.5 Concrete Policy Monitor	21
3 Control-Flow Integrity	23
3.1 Balancing between performance and security	23
3.1.1 Standard assumptions for effective <i>CFI</i>	24
3.2 Formal verification of Control-Flow Integrity	24
3.3 Control-Flow Integrity over PUMP	25

3.3.1	Enforcing Non-Writable Code & Non-Executable Data	25
3.3.2	Enforcing Control-Flow Integrity	26
	Bibliography	27

List of Figures

3.1	Rules enforcing NWC	25
3.2	Rules enforcing NXD	25

List of Listings

Chapter 1

Introduction

1.1 Motivation

Computer hardware and software continuously grow in size and complexity and as a result ensuring the absence of exploitable behaviors is becoming increasingly difficult. In the era when (**Feedback**: where?) computer systems are used extensively to carry important information (e.g. credit card numbers, national security documents), it has been widely accepted that security of these systems is a priority. Researchers have identified a number of potential vulnerabilities which arise from the violation of known but in-practice unenforceable safety and security policies.

So far, computer security has been delegated mostly to software, while the hardware is being almost completely controlled by the software. Programming languages have evolved, from low-level unmanaged languages, to high-level languages with features such as strong type systems and automatic memory management, making programming less error prone and reducing the number of exploitable bugs. Furthermore, in order to strengthen the security of computing systems a variety of mitigation techniques (**TODO**: reference some?) have been proposed, however these are mostly ad-hoc solutions designed to prevent specific known attacks, rather than enforcing a security policy along with a well defined class of attacks that are prevented, thus making it hard to reason about their effectiveness. In fact most of these mitigation techniques can be circumvented by attackers, (**TODO**: reference) which has lead to a continuous “chase” between attackers and security researchers.

One common attack technique is to deploy new code in the memory of the vulnerable program and then exploit some low-level vulnerability such as a buffer overflow to redirect the control flow to this attacker code. This attack can be stopped by a simple protection scheme known as $W \oplus X$, which enforces that a memory region is either executable or writable but not both. Unfortunately, new attack techniques can easily bypass $W \oplus X$. In particular, attackers have been using code-reuse attacks (e.g. return/jump - oriented programming) that allows them to chain together existing pieces of code to achieve malicious behavior without directly introducing new code.

The goal of this thesis is to describe and formalize in Coq a novel hardware-assisted implementation of an effective mitigation technique called Control Flow Integrity (CFI). CFI enforces that any execution of a program will respect a statically computed control flow graph (CFG), thus stopping a wide range of attacks that attempt to modify the control flow. As part of the formalization effort, following the work of Abadi *et al.* [1], we provide an attacker model and prove a variant of the CFI property described in [1].

1.2 Thesis Outline

Map 1. Intro 2a. Safety and Security Policies 2b. Micropolicies 3. CFI description 4. CFI formalization 5. Conclusions and Future work 6. Related work

Chapter 2 of this thesis briefly describes the basic requirements a security policy must satisfy and puts into context the framework we utilize in order to formalize and enforce a Control-Flow Integrity (CFI) policy. Chapter 3 discusses the current state of research on Control-Flow Integrity and clarifies our goals and contributions to it. Chapter 4 describes in detail the design of a fine-grained CFI policy and how we used the framework from Chapter 2 in order to enforce the policy and formally reason about its security properties. Chapter 5.. conclusions, future work? Chapter 6.. related work and bibliography? Appendix with code and/or step relations etc.?

Chapter 2

Safety and Security Policies

Currently the hardware provides only a small number of security mechanisms (**TODO**: name some), leaving most of the work to the software. This requires that the software performs various sanity-checks during an execution and that it carefully maintains various safety and security invariants, resulting in high runtime performance overheads.

Many potentially effective mitigation techniques are not deployed because of the performance overhead they incur. Another requirement for deployment of a protection mechanism is the compatibility with existing executables and the degree of intervention required by a human. Usually even making slight changes to a code (e.g. binary rewriting) and redistributing has high cost and the protection mechanism is likely to see very low adoption.

The lack of efficient and effective generic ways to enforce security policies, forces programmers to protect their own code, a task which is not trivial even for the small and simply programs. As a result most, if not all, software carries weaknesses which can be exploited by an attacker. “Safe” languages, automate some of the checks required and eases the work of the programmer, for example by implementing array bounds checking or by disallowing pointer-arithmetic. However these solutions only reduce the chance of introducing exploitable bugs in a program and do not enforce stricter, more effective policies such as Control Flow Integrity or complete Memory Safety (spatial/temporal protection for heap and stack). In addition, we still need effective and efficient protection mechanisms for a plethora of software written in unsafe languages such as C.

2.1 A Programmable Unit for Metadata Processing

The Programmable Unit for Metadata Processing (PUMP) architecture [6] allows us to efficiently implement a wide range of security policies [9] by associating metadata to the data being processed (e.g., this is an instruction, this is from the network, this is private), propagating the metadata as instructions are executed and using a rules-based system to check invariants on the metadata in parallel with the main computation. Abstractly, the tag propagation rules form a partial function from a set of input tags to a set of output tags

$$(opcode, tag_{pc}, tag_{instr}, tag_{arg1}, tag_{arg2}, tag_{arg3}) \mapsto (tag_{pc'}, tag_{result})$$

informally read as, “if the next instruction to be executed is *opcode*, the current tag of the program counter is *pc_{tag}*, the current tag on the instruction location is *tag_{instr}* and the tags on the operands of the instruction are *tag_{arg1}*, *tag_{arg2}* and *tag_{arg3}* then if execution

of the instruction is allowed the tag on the program counter should be set to tag_{pc} and any new data created by the instruction should be tagged tag_{result} ".

On the hardware level, the PUMP is an extension to a conventional RISC architecture. Every word of data in the machine - whether in the memory or a register, is extended with a word-sized metadata tag. These tags are not interpreted by hardware, instead the interpretation of the tags is left to the software, thus making it easy to implement new policies on the metadata. Since tags are word-sized, they can be pointers to complex data-structures of tags, such as tuples of tags, allowing for complex policies to be expressed and multiple orthogonal policies to be enforced in parallel.

The hardware undertakes the correct propagation of tags from operands to results according to the rules defined by the software. A hardware rule cache mapping sets of input tags to sets of output tags is used for common case efficiency. On each instruction dispatch, in parallel with the usual behavior of an instruction (e.g. execution of an addition in the ALU), the hardware forms the set of input tags and a lookup is performed on the rule cache. If the lookup is successful a set of output tags is returned and combined with the results of the normal execution of the instruction a new state is produced. On the other hand, if the lookup failed, the hardware invokes a trusted piece of system software - the fault handler - which checks the input tags and decides whether the execution should be allowed or not. In the first case, the fault handler returns a set of result tags, a pair of set of input and output tags is formed and inserted into the rules cache, while the faulting instruction is restarted and will now hit the cache. Otherwise, execution of this instruction violated some rules of the enforced policy and execution should not continue normally (e.g. should be halted).

As described in the original PUMP paper by Dehon *et al.* [6] and in more detail in the follow-up [9] a rich set of effective security policies can be efficiently implemented using the architecture mentioned above. In particular, implementations of dynamic typing, memory safety for heap-based data, control flow integrity and taint tracking are described, evaluated against a specific threat model and benchmarked. The benchmarks are done using a simulation of the described hardware and the two papers claim low overhead (10% on average) for each of the policies named above.

Compared to other software solutions for enforcing security policies, the PUMP offers significantly lower overhead, thanks to dedicated hardware assistance, while the fact that interpretation of the metadata is done by software offers flexibility with regard to the policies that can be implemented, compared to hardware solutions implementing a specific policy.

While the PUMP offers flexibility at a low runtime performance overhead, there are more overheads associated to such a mechanism. For example adding metadata to all the data in the machine, would result in a 100% memory overhead. In addition, the extra hardware and the rule cache along with potentially larger memories could result into a 400% overhead on energy usage. [9] The authors claim that a careful and well-optimized implementation can reduce these numbers, resulting in a 50% energy overhead.

2.2 Micro-policies: A Framework for Verified, Hardware-Assisted Security Monitors

The software components that can be changed to enforce a security policy are collectively called a micro-policy. Unsurprisingly, designing a security policy, reasoning about it's effectiveness against potential attackers and encoding it as a micro-policy can become a

complex task. Azevedo *et al.* [5] built a generic framework for defining micro-policies on top of a simple machine modeling a RISC processor augmented with the PUMP hardware (referred to as concrete machine), formalized this framework in Coq and used it to define and formally verify micro-policies for dynamic sealing, control-flow integrity, memory safety, compartmentalization and protecting the monitor code itself. (Feedback: maybe I should mention the word monitor at some point earlier)

The framework offers a high-level machine, called the symbolic machine, that abstracts away from unnecessary implementation details and can be used as an interface to the concrete machine, simplifying the work of the micro-policy designer. Additionally the symbolic machine is used to simplify correctness proofs. To instantiate the symbolic machine, the micro-policy designer needs to provide a set of symbolic tags which will be used to tag the various values of the machine, a transfer function that monitors program execution and determines how tags are propagated in each step and optionally a set of monitor services that are partial functions from machine states to machine states and can be used to control the monitor's behavior dynamically.

In order to implement the micro-policy at the concrete machine level, one needs to additionally provide machine code that implements the transfer function, an encoding of tags to words and machine code for any monitor services that the micro-policy may use. The relation between the symbolic and the concrete machine is formally defined as a two-way refinement (forward and backward). This is a generic refinement proof, parameterized by the encoding of the symbolic tags to words and a proof of correctness of the monitor code for a micro-policy. The designer of a micro-policy can use this two-way refinement simply by providing these two parameters.

2.2.1 Correctness of micro-policies

For each micro-policy an abstract machine which serves as a specification to the invariants the policy designer wants to enforce is defined. The abstract machine is “correct” by construction, meaning that it's designed to respect those invariants. Using the symbolic machine as an intermediate step to simplify the proofs, by proving a refinement between the symbolic and the abstract machine and by utilizing the generic refinement between the symbolic and the concrete machine, we can prove a refinement between the abstract and the concrete machine, thus showing that every valid step for the concrete machine is also a valid step for the abstract machine. (Feedback: say smth about steps and refinement earlier..)

2.2.2 Basic Machine

All the machines introduced in the original paper by Azevedo *et al.* [5], as well as this thesis, have a similar structure. In particular, they share a common RISC-based instruction set (with a few - uninteresting for the scope of this thesis - exceptions) and they have a fixed number of general-purpose registers, along with a pc register. Of course the abstract machine defined by the policy designer can differ in various ways, but more similarities with the symbolic machine implies easier proofs of correctness.

(TODO: write down a few rules?)

2.2.3 Symbolic Machine

As mentioned above, the symbolic machine enables us to abstract away from various low-level details of the concrete machine. We can express and reason about policies in terms of mathematical objects written in Gallina rather than machine code and the corresponding proofs for the concrete machine comes for free under some assumptions. The symbolic machine follows the structure of the basic machine but it's augmented to better match a PUMP architecture. Specifically the symbolic machine is parameterized by the following:

- A set of symbolic tags, used to tag the contents of the memory, the registers and the pc.
- A partial function *transfer*, that on every step checks whether the step is allowed according to opcode of the instruction executed and the tags on it. In the case it's allowed it returns a tag for the new pc and for any resulting data from executing the instruction.
- A partial function *get_service*, mapping addresses to *symbolic monitor services*. In the symbolic machine, monitor services are represented as a tuple of a partial function on machine states and a symbolic tag.
- An internal machine state with an initial value, that can be used by monitor services.

The states of the symbolic machine consists of a memory, registers, a *pc* register and an internal state. The memory and register contents, as well as the *pc*, are all tagged with a symbolic tag *t*. We name their contents *symbolic atoms* referred to with the notation $w@t$, where *w* is the value (word) and *t* is the tag.

At each step, a record named *mvector* is formed. It consists of the current opcode, the tag on the *pc*, the tag on the current instruction and optionally up to three tags depending on the opcode of the instruction. The *mvector* is passed to the transfer function which decides whether the step violated the policy enforced by the *transfer* function and in this case halts the machine, or if no violation occurred returns a tag for the new *pc* and a tag for any results the instruction execution produced.

(**TODO:** example rules, full rules in appendix?)

2.2.4 Concrete Machine

The concrete machine is a model of the basic machine with PUMP hardware, in particular a rules *cache* and a software *miss handler*. The instruction set has been extended with four additional instructions that are meant to be used by monitor code only, a restriction enforced by the monitor self-protection mechanism.

The states of the concrete machine consists of a memory, registers, a *pc* register, an *epc* register a special purpose register that holds the address of the faulting instruction after a cache miss and a cache. The cache works as a key-value store where a key is an *input vector* that contains an instruction opcode, the concrete tag of the current instruction, the concrete tag of the pc and up to three operand tags, and a value is an *output vector* which contain a tag for the new pc and a tag for any results from the execution of the instruction. Intuitively a concrete tag is the encoding into a word of a symbolic tag. Lifting this encoding relation to vectors, we get that a concrete vector is the encoding of a symbolic vector (*mvector*). In accordance (**Feedback:** em this sucks? does this word even exist? think about smth else?) to the symbolic machine the contents of the memory,

the registers, the pc and the epc are concrete atoms $w@t$ where w is a word and t is the encoding of a tag into a word.

The stepping relation for the concrete machine is a bit more complicated than the one for the symbolic machine. In particular, on each step the machine forms the *input vector* and looks it up in the cache. If the lookup succeeds then the instruction is allowed, a *output vector* is returned by the cache and the next state is tagged according to it. If the lookup fails, then the *input vector* is saved in memory, the current pc is stored in *epc* and the machine traps to the *miss handler*. The above are demonstrated in the two example rules below:

(**TODO:** put example rules)

Addresses 0 to 5 are used to store the *input vector* and 6 to 7 are used by the miss handler to store the *output vector*. As a side-note, cache eviction is not modeled (an infinite cache is assumed).

2.2.5 Concrete Policy Monitor

Unlike the symbolic machine, where the user cannot change the *transfer* function, enforcing a micro-policy on the concrete machine requires that we are able to protect the memory of the policy monitor and that privileged instructions are not executed by user code. This self-protection policy can be easily composed with another micro-policy and enforced by the infrastructure described above.

Using tags of the form, *User st*, *Entry st*, *Monitor* we can distinguish between user memory, monitor memory and monitor services. In particular *User st* is used to tag a user-level atom, where *st* is the word-encoding of a symbolic tag. *Monitor* is used to tag the monitor memory and a few reserved registers. The pc is tagged with *Monitor* when a monitor execution takes place and *User st* when user-code is executed. The tag *Entry st* is used to tag the first instruction of a monitor service and serves as an indication that execution will continue under the privileged *Monitor* mode.

One important thing to note is that the miss handler for the concrete machine can take an arbitrary number of steps before deciding that no violation occurred and returning to *User* mode, unlike the symbolic *transfer* function that does not need to take any steps.

Chapter 3

Control-Flow Integrity

Restricting the control-flow of a program in some way is a technique widely spread among security researchers. For example non-executable data (NXD) can be considered as a form of (very) coarse-grained *CFI* where control-flow is not allowed to reach any memory region that holds non-executable data. Other mitigation techniques such as protecting return addresses on the stack enforce a form of coarse-grained *CFI*.

Moreover it is common that security properties are enforced dynamically by code that is statically injected to the program (e.g., Inlined Reference Monitors (IRM) [7] follow that approach), thus some form of *CFI* is required in order to ensure that these checks are not circumvented.

(**TODO:** Think about title)

3.1 Balancing between performance and security

Abadi *et al.* first proposed a technique to enforce *CFI* based on IRMs. In particular, they proposed to mark all valid targets of *indirect* control transfers with a unique identifier and inject checks before all indirect jumps (including return instructions). However they assume that any two destinations are equivalent, in the sense that they share the same identifier, if the CFG contains edges from the same set of sources, which may significantly reduce the precision of the CFG. The authors also note that a 2-ID approach where one identifier is used for calls and another for returns could provide adequate security in many cases.

The work of Abadi *et al.* sparked interest of researchers who tried to improve some of the weaknesses of the initial implementation, usually by choosing between performance against precision and vice-versa.

Bletsch *et al.* [3] followed the work of Abadi *et al.*, but changed their checking mechanism to perform the check after the control flow transfer has occurred which, as the authors claim, reduced the cache pressure and resulted in better performance. Precision remains the same with the implementation of Abadi *et al.*.

Zhang *et al.* [11] proposed *Compact Control Flow Integrity and Randomization* (CCFIR), a new efficient way to enforce coarse-grained *CFI*. CCFIR collects all valid targets of indirect control-transfers and stores them in a random order, in a protected section called “Springboard section”. Indirect control-transfers are only allowed to addresses that are in the Springboard. Their implementation uses a 3-ID approach where one identifier is used for calls and the two other identifiers are for returns, separating them between returns to sensitive and non-sensitive functions. Their implementation also supports interaction

between protected and un-protected modules, which makes it an attractive solution to coarse-grained *CFI*.

The above techniques are evaluated in [8] where the authors demonstrate code-reuse attacks against binaries protected by coarse-grained *CFI*. These attacks illustrate the need for fine-grained *CFI* which however incurs a high runtime-overhead penalty making deployment of such a mechanism unlikely.

3.1.1 Standard assumptions for effective *CFI*

Most -if not all- *CFI* implementations also come with a set of assumptions under which *CFI* holds. Two standard assumptions for all mechanisms that attempt to enforce *CFI* are:

- *NXD* is an abbreviation for Non-Executable Data, a security mechanism that disallows execution of data.
- *NWC* stands for Non-Writable Code. Changing the code of a program would allow an attacker to circumvent dynamic checks.

Both assumptions are fairly standard for modern computers and are enforced through hardware or software. In some cases *NXD* can be lifted, but additional security risks and complexity is not worth the minor advantages offered by such an action.

Many implementations that attempt to do fine-grained *CFI* also require that identifiers used to mark nodes in the CFG are unique.

3.2 Formal verification of Control-Flow Integrity

In [2] Abadi *et al.* extended their original paper, with -among other things- a more detailed formal study of *CFI*. Their formalization regarded a much simpler machine than the x86 omitting all the complexity in modern systems. The machine has a few instructions, a separate data memory and instruction memory which by the operational semantics of the machine are non-executable and non-writable (enforcing *NXD* and *NWC* by construction), and a small set of registers. Moreover, their attacker model permits arbitrary changes to the data memory, arbitrary changes to all the registers but a few distinguished ones that are used during the dynamic checks and no changes to the instruction memory. The authors prove that under some assumptions *CFI* is preserved for every step even in the presence of an attacker as powerful as the one described above. Their formal study served as a guideline for the implementation, but as it is done on paper their proofs cannot be machine checked. Furthermore, their formalization omits less interesting but important details such as instruction encoding and decoding which as shown in [10] are far from trivial for the x86.

Machine-checked formal verification efforts include [12], which is a SFI formalization for the ARM architecture that also enforces *CFI*. Their formalization was developed using the HOL theorem prover and a program logic framework they created. However their benchmarks report a 240% runtime overhead. The authors of [4] claim partial proofs for a *CFI* enforcement mechanism focused on the kernel of an operating system. Their runtime overhead can also reach 100%.

3.3 Control-Flow Integrity over PUMP

The PUMP hardware allows us to avoid taking the difficult decision between performance and security. As shown in [9], we can enforce a *fine-grained CFI* policy with an average overhead of 8%. (TODO: Is this number right?)

In our design, we take the standard approach and claim *CFI* under *NXD* and *NWC*. We considered designs that lifted these assumptions but we rejected them, for the time being, as there did not seem to be any considerable advantage i.e., compatibility with self-modifying programs, JIT compilers, etc. Allowing the code of the program to change, would in practice require for the CFG to change as well, which unless done in a controlled, “safe” way, would invalidate the enforcement of *CFI*. However, we do not have to rely on special hardware or software to enforce *NXD* and *NWC*. We can achieve this easily and efficiently by creating a separate micro-policy.

3.3.1 Enforcing Non-Writable Code & Non-Executable Data

Consider the set of tags $\mathcal{T} = \{Data, Instr\}$. If we initially tag all executable regions in memory as *Instr* and all non-executable as *Data* then we can enforce *NWC* by two rules of the form

$$\begin{aligned} (Store, _, _, _, _, Instr) &\rightarrow \emptyset \\ (Store, _, _, _, _, Data) &\rightarrow (_, Data) \end{aligned}$$

Figure 3.1: Rules enforcing *NWC*

The $_$ in the vectors, represent *don't care* values. In the context of the input vector their behavior is the same as *don't care* values in match expressions in ML languages. In the context of the output vector it just captures the intuition that we will not really use the result tags, so anything could be returned as a result tag (i.e., *Data* or we can copy-through tags from the input vector). Informally the above rules reads as “If the current opcode is *Store* and the content of the memory location we are trying to write is tagged *Instr* then the memory write is not allowed. Otherwise if it is tagged *Data* then the write is permitted and the new value will also be tagged *Data*.”

We can enforce *NXD* in a similar fashion

$$\begin{aligned} (-, -, Data, -, -, -) &\rightarrow \emptyset \\ (-, -, Instr, -, -, -) &\rightarrow (-, -) \end{aligned}$$

Figure 3.2: Rules enforcing *NXD*

Informally the above rules reads as “If the tag on the current instruction is *Data* then execution is not allowed. Otherwise if it is *Instr* then execution is allowed”.

(Feedback: Used $_$ and $-$, I think the second one looks better, opinions?)

(TODO: Perhaps explain what each tag means for each opcode earlier – or maybe just in appendix?)

(**Feedback:** These tuple-vectors make it hard for people not familiar with them to remember what each field is, any better ways to represent them?)

3.3.2 Enforcing Control-Flow Integrity

Bibliography

- [1] M. Abadi, M. Budiu, Ú. Erlingsson, and J. Ligatti. Control-flow integrity. In *12th ACM Conference on Computer and Communications Security*, pages 340–353. ACM, 2005.
- [2] M. Abadi, M. Budiu, Ú. Erlingsson, and J. Ligatti. Control-flow integrity principles, implementations, and applications. *ACM Transactions on Information System Security*, 13(1), 2009.
- [3] T. Bletsch, X. Jiang, and V. Freeh. Mitigating code-reuse attacks with control-flow locking. In *Proceedings of the 27th Annual Computer Security Applications Conference, ACSAC '11*, pages 353–362, New York, NY, USA, 2011. ACM.
- [4] J. Criswell, N. Dautenhahn, and V. Adve. KCoFI: Complete control-flow integrity for commodity operating system kernels. 2014.
- [5] A. A. de Amorim, M. Dénès, N. Giannarakis, C. Hrițcu, B. C. Pierce, A. Spector-Zabusky, and A. Tolmach. Micro-policies: A framework for verified, hardware-assisted security monitors. Under Review, July, July 2014.
- [6] U. Dhawan, N. Vasilakis, R. Rubin, S. Chiricescu, J. M. Smith, T. F. Knight, B. C. Pierce, and A. DeHon. PUMP – A Programmable Unit for Metadata Processing. In *Proceedings of the 3rd International Workshop on Hardware and Architectural Support for Security and Privacy, HASP '14*, New York, NY, USA, June 2014. ACM.
- [7] Ú. Erlingsson. *The inlined reference monitor approach to security policy enforcement*. PhD thesis, Cornell University, Jan. 2004.
- [8] E. Göktas, E. Athanasopoulos, H. Bos, and G. Portokalidis. Out of control: Overcoming control-flow integrity. In *IEEE Symposium on Security and Privacy*, 2014.
- [9] C. Hrițcu, U. Dhawan, N. Vasilakis, S. Chiricescu, J. M. Smith, B. C. Pierce, and A. DeHon. Programming the PUMP: Hardware-assisted micro-policies for security. Under Review, May, May 2014.
- [10] G. Morrisett, G. Tan, J. Tassarotti, J.-B. Tristan, and E. Gan. RockSalt: better, faster, stronger SFI for the x86. In *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*, pages 395–404. ACM, 2012.
- [11] C. Zhang, T. Wei, Z. Chen, L. Duan, L. Szekeres, S. McCamant, D. Song, and W. Zou. Practical Control Flow Integrity & Randomization for Binary Executables. In *IEEE Symposium on Security and Privacy*, 2013.

- [12] L. Zhao, G. Li, B. D. Sutter, and J. Regehr. ARMor: fully verified software fault isolation. In *11th International Conference on Embedded Software*, pages 289–298. ACM, 2011.